

Introduction To HPCC

February 27, 2014

Dirk Colbry

colbrydi@msu.edu

Director, High Performance Computing Center

Institute for Cyber-Enabled Research

MICHIGAN STATE
UNIVERSITY

© 2014 Michigan State University Board of Trustees.




Agenda

- Short introduction of HPCC and iCER
- Quick overview of the HPCC System and submission scripts
- Review of SampleQFC.qsub
- Questions/Discussion of lab's workflow and possible improvements

MICHIGAN STATE
UNIVERSITY






2004 HPCC

- Provide a level of performance beyond what you could get and reasonably maintain as a small group
- Provide a variety of technology, hardware and software, that would allow for innovation not easily found

MICHIGAN STATE
UNIVERSITY

2009 iCER

The Institute for Cyber Enabled Research(iCER) at Michigan State University (MSU) was established to coordinate and support multidisciplinary resource for computation and computational sciences. The Center's goal is to enhance MSU's national and international presence and competitive edge in disciplines and research thrusts that rely on advanced computing.

MICHIGAN STATE
UNIVERSITY



HPC Systems

FREE*


- Large Memory Nodes (up to 2TB!)
- GPU Accelerated cluster (K20, M1060)
- PHI Accelerated cluster (5110p)
- Over 540 nodes, 10000 computing cores
- Access to high throughput condor cluster
- 363TB high speed parallel scratch file space
- 50GB replicated file spaces
- Access to large open-source software stack and specialized bioinformatics VMs

MICHIGAN STATE
UNIVERSITY



Free Access to software

- Compiled open-source software stack
 - over 300 titles!
- Optimized Math/Communications libraries
- Some commercial software available
 - E.g. Ansys, MATLAB (+many toolboxes), Stata, Gauss, SAS



MICHIGAN STATE
UNIVERSITY




What if I want more?

XSEDE

Extreme Science and Engineering
Discovery Environment



Open Science Grid



MICHIGAN STATE
UNIVERSITY



What if I need help?

- Ask us!
- Local Workshops
 - Software carpentry
 - Introduction to Linux and HPCC
 - Advanced HPCC
- Remote Training
 - VSCSE – Virtual School for Computer Science Education
 - XSEDE training Workshops

 software carpentry



XSEDE

Extreme Science and Engineering
Discovery Environment



MICHIGAN STATE
UNIVERSITY




MSU Seminars in Research and Instructional Technology

May 6,7 2014

- Two days of no-cost seminars to faculty and graduate students on technology topics.
 - Morning sessions run from 8:30 to 11:30 am.
 - Afternoon sessions run from 1:30 to 4:30 pm.
 - Lunch is provided that will feature guest speakers on instructional technology.
- Tuesday May 6
 - Introduction to HPC
 - Advanced HPC
- Wednesday May 7
 - Managing, Sharing and Moving Big Data

MICHIGAN STATE
UNIVERSITY

<http://train.msu.edu/faculty/seminars/>



Submission Script

1. List of required resources
2. All command line instructions needed to run the computation

Typical Submission Script

```
#!/bin/bash -login
#PBS -l walltime=10:00:00,mem=3Gb,nodes=10:ppn=1
#PBS -j oe

cd ${PBS_O_WORKDIR}

./myprogram -my input arguments

qstat -f ${PBS_JOBID}
```

Shell Comment

Define Shell

Resource Requests

Shell Commands

Special Environment Variables

simple.qsub


```
#!/bin/bash -login
#PBS -l walltime=00:01:00
#PBS -l nodes=1:ppn=1,feature=gbe

cd ${PBS_O_WORKDIR}
./hello

qstat -f ${PBS_JOBID}
```

Submitting a job

- `qsub -arguments <Submission Script>`
 - Returns the job ID. Typically looks like the following:
 - 5945571.cmgr01
- Time to job completion



System Limitations

- Scheduling
 - 10 eligible jobs at a time
 - 384 running jobs
- Resources
 - 1 week of walltime
 - 384 cores (nodes * ppn)
 - ppn=64
 - 2TB memory on a single core
 - ~200 GB Hard Drive

Scheduling Priorities

- Jobs that use more resources get higher priority (because these are hard to schedule)
- Smaller jobs are backfilled to fit in the holes created by the bigger jobs
- Eligible jobs acquire more priority as they sit in the queue
- Jobs can be in three basic states:
 - Blocked, eligible or running



Job completion

- By default the job will automatically generate two files when it completes:
 - Standard Output:
 - Ex: jobname.o5945571
 - Standard Error:
 - Ex: jobname.e5945571
- You can combine these files if you add the join option in your submission script:
 - “#PBS -j oe”



One-on-one consulting

SampleQFC.qsub

First Look

```
#!/bin/sh -login
^M
^M
# Time job will take to execute (HH:MM:SS format)^M

#PBS -l walltime=48:00:00
^M
^M
# Memory needed by the job^M

#PBS -l mem=5Gb      ^M

^M
```

dos2unix

- Converts special end-of-line characters from windows format to dos.

dos2unix SampleQFC.qsub

- You could also use an editor that uses unix end-of-line characters.
 - I recommend notepad++

Extra Errors:

```
# Print out values of the current jobs PBS environment
variables

++++++

#./nbhurygb -l1 50000000 -l2 50000000 -l3 10000000 -n11
10000000 -n13 10000000

#Convert files to unix format
dos2unix *.*

#here call your bash file or your executable file

++++++

admb -r nbhurygb
```

Three Types of Comments

```
#!/bin/sh -login

# Time job will take to execute (HH:MM:SS format)
#PBS -l walltime=48:00:00

# Memory needed by the job
#PBS -l mem=5Gb

# Number of shared memory nodes required and the number
of processors per node
#PBS -l nodes=1:ppn=1,feature=gbe

# Make output and error files the same file
#PBS -j oe
```

Short 1/2

```
#!/bin/sh -login
#PBS -l walltime=48:00:00
#PBS -l mem=5Gb
#PBS -l nodes=1:ppn=1
#PBS -j oe
#PBS -m abe
#PBS -N nbhurygb
#PBS -l file=5gb

# Change to the Original Working Directory
cd ${PBS_O_WORKDIR}
cd nbhurygb

# Print PBS Environment Variables
env | grep PBS
```

Short 2/2

```
#Convert files to unix format
dos2unix *.*

#here call your bash file or your executable file
admb -r nbhurygb

#Send output to this file
./nbhurygb >runtime.log

qstat -f ${PBS_JOBID}
```

Example Modifications

- Use feature=gbe
 - Allows you to run on a few extra nodes
- Use your JobName to make script more portable.

Original 1/2

```
#!/bin/sh -login
#PBS -l walltime=48:00:00
#PBS -l mem=5Gb
#PBS -l nodes=1:ppn=1
#PBS -j oe
#PBS -m abe
#PBS -N nbhurygb
#PBS -l file=5gb

# Change to the Original Working Directory
cd ${PBS_O_WORKDIR}
cd nbhurygb

# Output Contents of the PBS NODEFILE
env | grep PBS
```

Modified 1/2

```
#!/bin/sh -login
#PBS -l walltime=48:00:00
#PBS -l mem=5Gb
#PBS -l nodes=1:ppn=1,feature=gbe
#PBS -j oe
#PBS -m abe
#PBS -N nbhurygb
#PBS -l file=5gb

# Change to the Original Working Directory
cd ${PBS_O_WORKDIR}
cd ${PBS_JOBNAME}

# Output Contents of the PBS NODEFILE
env | grep PBS
```

Original 2/2

```
#Convert files to unix format
dos2unix *.*

#here call your bash file or your executable file
admb -r nbhurygb

#Send output to this file
./nbhurygb >runtime.log

qstat -f ${PBS_JOBID}
```

Modified 2/2

```
#Convert files to unix format
dos2unix *.*

#here call your bash file or your executable file
admb -r ${PBS_JOBNAME}

#Send output to this file
./${PBS_JOBNAME} > runtime.log

qstat -f ${PBS_JOBID}
```

Command line option

- Now instead of a different script for each job you can just use one script:

```
qsub -N nbhurygb_2 SampleQFC.qsub
```

Requesting local disk

- Sometimes (not often) local disk is faster than scratch
- Users can use the following resource to request temporary local disk space:
 - #PBS -l file=10gb
- The directory to access this disk space is determined by the one time use environment variable
 - \${TMPDIR}

1/2

```
#!/bin/sh -login
#PBS -l walltime=48:00:00
#PBS -l mem=5Gb
#PBS -l nodes=1:ppn=1,feature=gbe
#PBS -j oe
#PBS -m abe
#PBS -N nbhurygb
#PBS -l file=5gb

# Change to the Original Working Directory
cd ${PBS_O_WORKDIR}
cd ${PBS_JOBNAME}

# Output Contents of the PBS NODEFILE
env | grep PBS
```


2/2

```
#Convert files to unix format
dos2unix *.*

#here call your bash file or your executable file
admb -r ${PBS_JOBNAME}

#Send output to this file
./${PBS_JOBNAME} > runtime.log

qstat -f ${PBS_JOBID}
```

Discussion